

Hypothesis Testing

Part 1: One and two sample t-tests

QuantArch Week 5 | 07-03-2022



Universiteit
Leiden

One sample t test

$$t = \frac{\bar{x} - \mu}{s_x} \quad s_x = \frac{s}{\sqrt{n}}$$

One sample t test

Let's test the **Area** variable against a population mean.

Wait... have you uploaded your data?

```
pits_data <- readr::read_csv(here("data/derived_data/pits-data.csv"))
```

One sample t test

Let's test the **Area** variable against a population mean.

We will randomly generate the true population mean μ from a normal distribution.

```
set.seed(42) # make sure we generate the same number every time
(pop_mean <- rnorm(1, mean = mean(pits_data$Area, na.rm = T), 1))
```

```
## [1] 238.9387
```

The mean of Area is `mean(pits_data$Area, na.rm = T)` = **237.5677778**

The standard deviation is `sd(pits_data$Area, na.rm = T)` = **94.4995673**

The sample size is `length(!is.na(pits_data$Area))` = **91**


One sample t-test

$$t = \frac{\bar{x} - \mu}{s_x} \quad s_x = \frac{s}{\sqrt{n}}$$

Now we can just run that through our calculator...

```
t_stat <- (237.568 - 238.939) / (94.500 / sqrt(91))
```

Our t-statistic is **-0.1383969**



$$t = \frac{\bar{x} - \mu}{s_x}$$

One sample t-test: A better option

We can translate the t-test equation to code

First the nominator (or difference in means)

```
t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean
```

then the sample standard deviation

```
sd_sample <- sd(pits_data$Area, na.rm = T)
```

and the sample size... why **sum** not **length**?

```
n_sample <- sum(!is.na(pits_data$Area))
```

Put it together

```
t_stat <- t_nom / (sd_sample / sqrt(n_sample))
```

Our t-statistic is **-0.1376308**.

Now how do we evaluate the t-statistic?

We can calculate the critical t-value

First we set out α . Let's say 0.047212... it's just as arbitrary as 0.05


```
alpha <- 0.047212
```

Now we use the alpha and the degrees of freedom (df) to calculate the critical t-value

```
qt(1 - alpha, n_sample - 1)
```

```
## [1] 1.690541
```

Let's see where our t-value falls within the t-distribution with the calculate critical t-values.



Can we use this method to calculate t for other variables, too?



One sample t-test: A betterer option

Functions

We can create a function!

Starting with a name

```
t_stat <- function(){  
  # body of the function  
}
```

Then we can take the code we used earlier,

```
t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean  
sd_sample <- sd(pits_data$Area, na.rm = T)  
n_sample <- sum(!is.na(pits_data$Area))  
t_stat <- t_nom / (sd_sample / sqrt(n_sample))
```

One sample t-test: A betterer option

Functions

We can create a function!

Starting with a name

```
t_stat <- function(){  
  # body of the function  
}
```

Then we can take the code we used earlier, and put it inside the body of the function

```
t_calc <- function(){  
  t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean  
  sd_sample <- sd(pits_data$Area, na.rm = T)  
  n_sample <- sum(!is.na(pits_data$Area))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
}
```

One sample t-test: A betterer option

Functions

and now we just need the value we want the function to **return**.

```
t_calc <- function() {  
  t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean  
  sd_sample <- sd(pits_data$Area, na.rm = T)  
  n_sample <- sum(!is.na(pits_data$Area))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
}
```

One sample t-test: A betterer option

Functions

we need to specify which object we want the function to **return**.

```
t_calc <- function(){  
  t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean  
  sd_sample <- sd(pits_data$Area, na.rm = T)  
  n_sample <- sum(!is.na(pits_data$Area))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
  return(t_stat)  
}
```

We can run the function by calling its name

```
t_calc()
```

```
## [1] -0.1376308
```

One sample t-test: A betterer option

Functions

To make sure we can use the same function for different variables, we need to give it arguments and remove **Area**-specific objects

```
t_calc <- function(){  
  t_nom <- mean(pits_data$Area, na.rm = T) - pop_mean  
  sd_sample <- sd(pits_data$Area, na.rm = T)  
  n_sample <- sum(!is.na(pits_data$Area))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
  return(t_stat)  
}
```


One sample t-test: A betterer option

Functions

To make sure we can use the same function for different variables, we need to give it arguments and remove **Area**-specific objects

```
t_calc <- function(variable, pop_mean){  
  t_nom <- mean(variable, na.rm = T) - pop_mean  
  sd_sample <- sd(pits_data$Area, na.rm = T)  
  n_sample <- sum(!is.na(pits_data$Area))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
  return(t_stat)  
}
```

One sample t-test: A betterer option

Functions

To make sure we can use the same function for different variables, we need to give it arguments and remove **Area**-specific objects

```
t_calc <- function(variable, pop_mean){  
  t_nom <- mean(variable, na.rm = T) - pop_mean  
  sd_sample <- sd(variable, na.rm = T)  
  n_sample <- sum(!is.na(variable))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
  return(t_stat)  
}
```

One sample t-test: A betterer option

Functions

To make sure we can use the same function for different variables, we need to give it arguments and remove **Area**-specific objects

```
t_calc <- function(variable, pop_mean){  
  t_nom <- mean(variable, na.rm = T) - pop_mean  
  sd_sample <- sd(variable, na.rm = T)  
  n_sample <- sum(!is.na(variable))  
  t_stat <- t_nom / (sd_sample / sqrt(n_sample))  
  return(t_stat)  
}
```

The last two lines don't need changing, because all the objects are calculated within the function.

Exercise

Use the function to calculate the t-statistic and critical t-value for **Length**

We'll pretend that the true population mean is **1.72**, and we'll set our α at **0.05**.

Solution

```
# t-statistic
t_stat <- t_calc(pits_data$Points, 1.72)
t_stat
```

```
## [1] 2.159668
```

```
# p-value
alpha <- 0.05
df <- sum(!is.na(pits_data$Points)) - 1
qt(1 - alpha, df)
```


```
## [1] 1.662354
```

Break

Evaluating the t-statistic

Let's look at the critical values for our **Points** statistic.

As you may be able to tell, the overall process is not a dichotomy.



Evaluating the t-statistic

As you may be able to tell, the overall process is not a dichotomy.

We can produce a probability value of the data/estimate in relation to the 'null hypothesis'.

p-value

Informally, a p-value is the probability under a specified statistical model that a statistical summary of the data (e.g., the sample mean difference between two compared groups) would be equal to or more extreme than its observed value. (American Statistical Association)

The p-value is NOT the probability of your hypothesis. Or any hypothesis for that matter...

Evaluating the t-statistic

Our α determines the critical threshold for our p-value in order to reject the null hypothesis.

Our α from the previous example is 0.05,

which means we need a p-value equal to, or lower than 0.05.

Calculating the p-value

```
2 * pt(-abs(t_stat), n_sample - 1)
```

```
## [1] 0.03348712
```


Two sample t-test

Exercise

Create a function to calculate the t-statistic of a two sample t-test

$$t = \frac{\bar{x} - \bar{y}}{s_{x-y}}$$

$$s_{x-y} = \sqrt{\frac{(n_x - 1)s_x^2 + (n_y - 1)s_y^2}{n_x + n_y - 2} \left(\frac{1}{n_x} + \frac{1}{n_y} \right)}$$

💡 **Hint:** Code > Rainbow Parentheses

Solution

```
t_calc2 <- function(x, y){  
  mean_x <- mean(x, na.rm = T)  
  mean_y <- mean(y, na.rm = T)  
  n_x <- sum(!is.na(x))  
  n_y <- sum(!is.na(y))  
  s_x <- sd(x, na.rm = T)  
  s_y <- sd(y, na.rm = T)  
  
  s_xy <- sqrt(  
    (((n_x-1)*s_x^2 + (n_y-1)*s_y^2) /  
      (n_x + n_y -2)) *  
      (n_x^-1 + n_y^-1)  
    )  
  t_stat <- (mean_x - mean_y) / s_xy  
  return(t_stat)  
}
```

Two sample t-test

Now let's compare Areas of each **Location**.

```
inside <- filter(pits_data, Location == "Inside")
outside <- filter(pits_data, Location == "Outside")
t_stat2 <- t_calc2(inside$Area, outside$Area)
t_stat2
```

```
## [1] 9.901939
```

```
df <- sum(!is.na(pits_data$Area)) - 1
p_value <- 2 * pt(-abs(t_stat2), df = df)
p_value
```

```
## [1] 5.104513e-16
```

We can check our results from the one and two sample t-tests with the built-in function `t.test`

For a one sample test we need to specify `mu`

```
pop_mean <- 238.9387  
t.test(pits_data$Area, mu = pop_mean)
```

```
##  
##      One Sample t-test  
##  
## data:  pits_data$Area  
## t = -0.13763, df = 89, p-value = 0.8908  
## alternative hypothesis: true mean is not equal to 238.9387  
## 95 percent confidence interval:  
##  217.7752 257.3603  
## sample estimates:  
## mean of x  
##  237.5678
```

For a two sample test we need the two variables we are comparing (and `var.equal = T` for now).

```
t.test(inside$Area, outside$Area, var.equal = T)
```

```
##  
##      Two Sample t-test  
##  
## data:  inside$Area and outside$Area  
## t = 9.9019, df = 88, p-value = 5.715e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  110.4006 165.8415  
## sample estimates:  
## mean of x mean of y  
##  317.3711  179.2500
```



Yes, we could have done that all along...

But you learned something along the way, right? ...right??

Reporting results

The mean area of house pits Inside (m = 317) the "white wall" is 138 sqft larger than the mean area of houses outside (m = 179) the "white wall", $t(89) = 9.9$, $p < 0.001$, 95%CI(110, 166).

Note: Report exact p-value if $p > 0.001$

Or a table

```
t.test(inside$Area, outside$Area, var.equal = T) %>%  
  tidy()
```

```
## # A tibble: 1 × 10  
##   estimate estimate1 estimate2 statistic  p.value parameter conf.low conf.high  
##   <dbl>     <dbl>     <dbl>     <dbl>    <dbl>     <dbl>     <dbl>     <dbl>  
## 1    138.     317.     179.     9.90 5.72e-16         88    110.     166.  
## # ... with 2 more variables: method <chr>, alternative <chr>
```